
Pretraining and Finetuning LLMs
from the Ground Up

	Workshop topics
1	Introduction to LLMs
2	Understanding LLM input data
3	Coding an LLM architecture
4	Pretraining LLMs
5	Loading pretrained weights
6	Finetuning LLMs

Developing an LLM

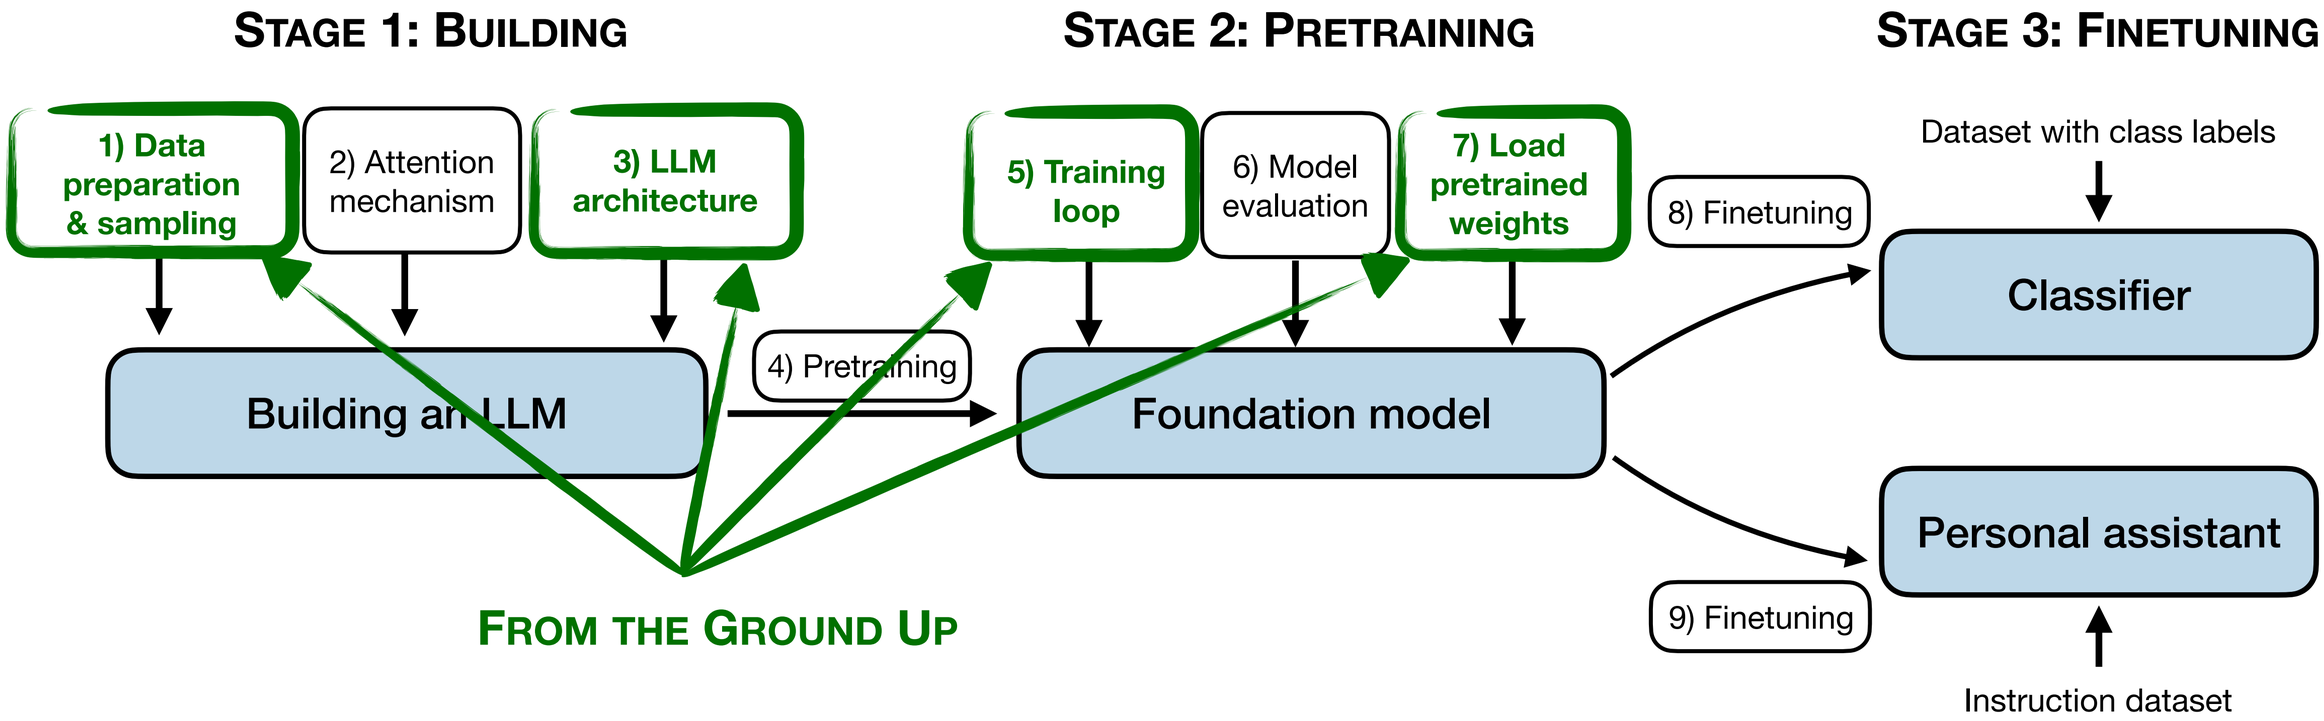

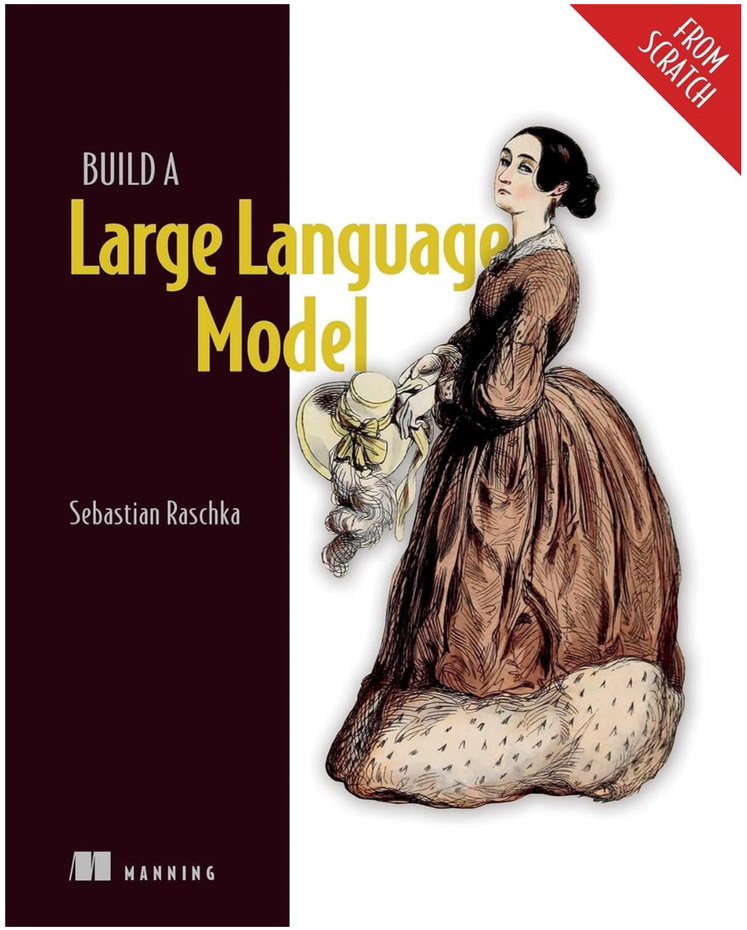

https://mng.bz/lrp2

https://github.com/rasbt/LLMs-from-scratch

(Source for most figures and code)

Developing an LLM

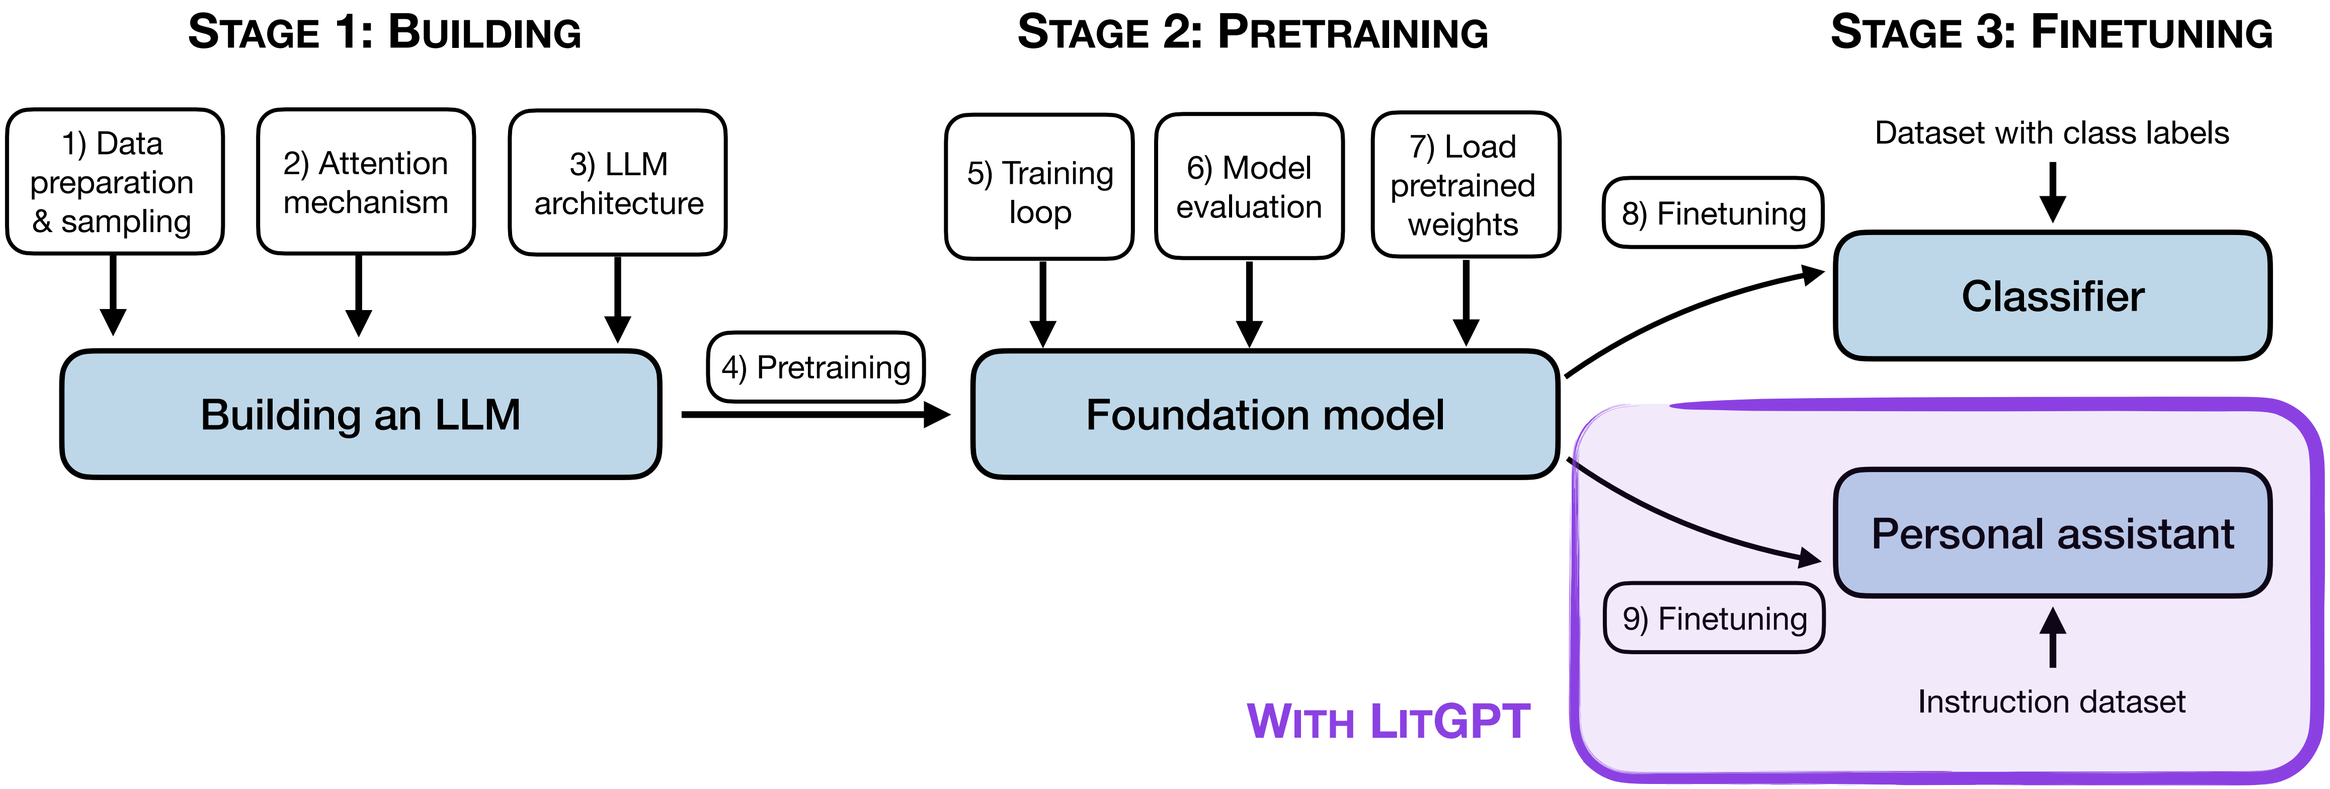

README License

⚡ LitGPT

20+ high-performance LLM implementations with recipes to pretrain, finetune, deploy at scale.

✅ From scratch implementations
✅ Flash attention
✅ Reduce GPU memory (fp4/8/16/32)
✅ No abstractions
✅ FSDP
✅ 1-1000+ GPUs/TPUs
✅ Beginner friendly
✅ LoRA, QLoRA, Adapter
✅ 20+ LLMs

Lightning AI • Quick start • Models • Finetune • Deploy • All workflows • Features • Recipes (YAML) • Tutorials

Get started

https://github.com/Lightning-AI/litgpt

Lightning AI

Creators of PyTorch Lightning

Simple. Powerful.

Zero setup. Persistent. Always ready.

Studio marries the simplicity of a **local development experience** with the power of **1,000s of cloud GPUs**, unlimited storage and multiplayer collaboration.

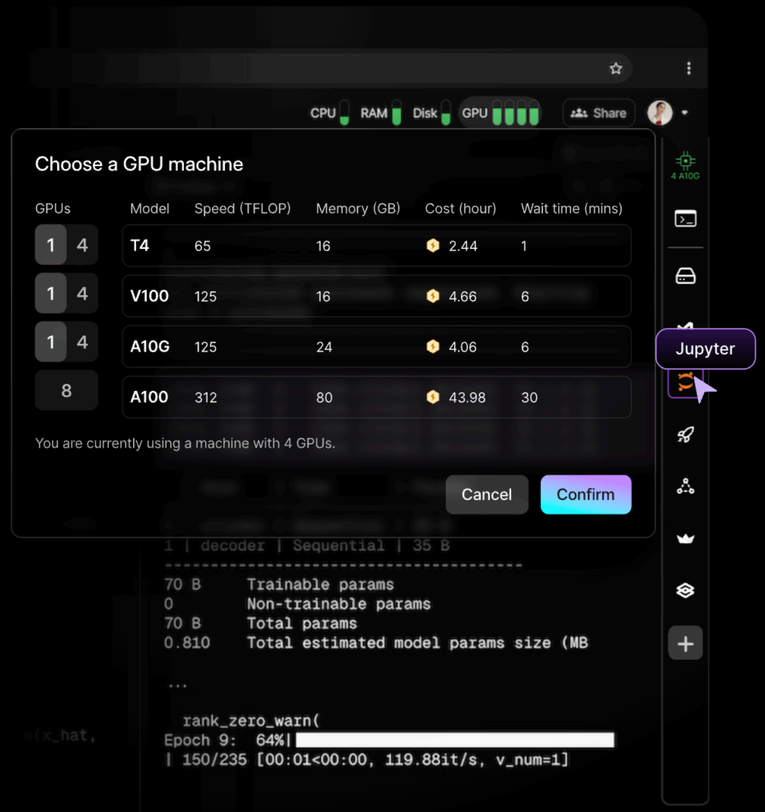

- No environment setup.
- Code in the browser or connect your local IDE.
- **Switch from CPU to GPU with zero environment changes.**
- Host and share AI apps. Streamlit. Gradio. React JS.
- Code together.
- Infinite storage. Upload, share files and connect S3 buckets.

https://lightning.ai/

Source

Lightning AI Public

Explore

Featured
Trending
Recent
All studios
My studios

Educational

Blogs
Papers
Tutorials

Workflows

Data processing
Endpoints
Training
Serving
Other

Model types

Audio
Image
Multimodal
Text
Tabular

Document Chat Assistant using RAG

aniket 269 6.54 K

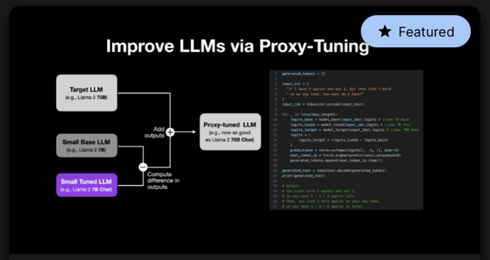

Improve LLMs With Proxy-Tuning

sebastian 47 7.88 K

Embed English Wikipedia under 5 dollars

thomasgridai 26 2.73 K

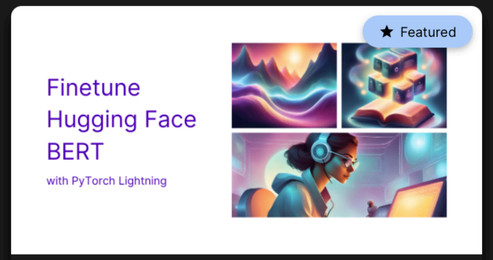

Finetune Hugging Face BERT with PyTorch Lig...

JG justin 97 1.98 K

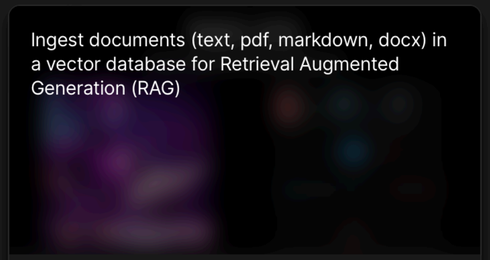

Document Search and Retrieval using RAG

aniket 676 7.10 K

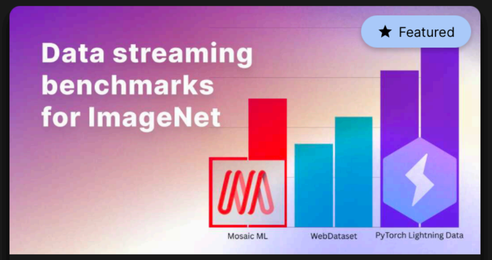

Benchmark cloud data-loading libraries

thomasgridai 23 1.05 K

Prepare the TinyLlama 1T token dataset

thomasgridai 38 1.64 K

Code LoRA from Scratch

sebastian 229 24.66 K

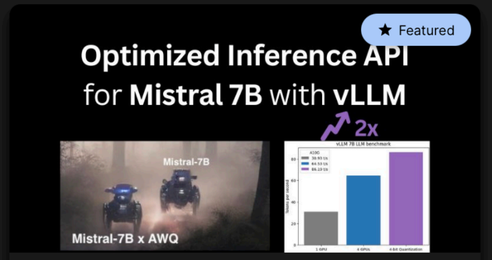

Optimized LLM inference API for Mistral 7B usi...

aniket 50 7.63 K

Contact

 @rasbt  in/sebastianraschka

 https://sebastianraschka.com/contact/

 https://lightning.ai